Utility Component Scores Across Experimental Settings

Absolute scores · all tested weights · three-seed means for trained configurations

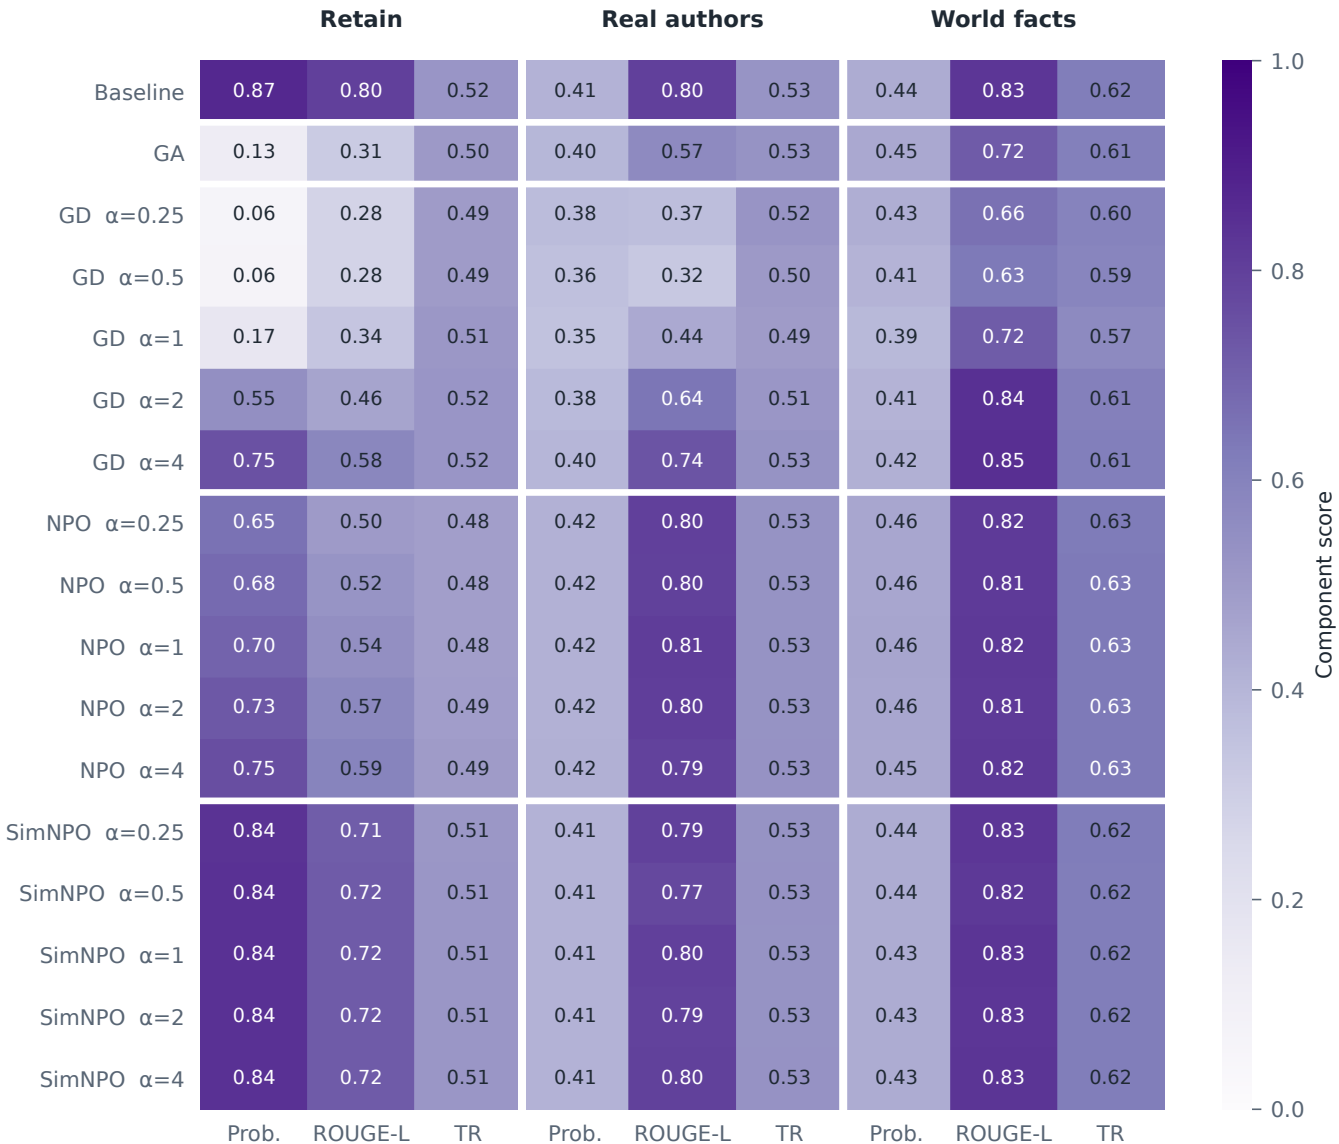

Supplementary figure S1. Baseline is a fixed reference (n = 1); each trained setting has n = 3 seeds.
Prob. = answer probability; ROUGE-L = recall; TR = truth ratio. Darker purple denotes a higher score (0–1).