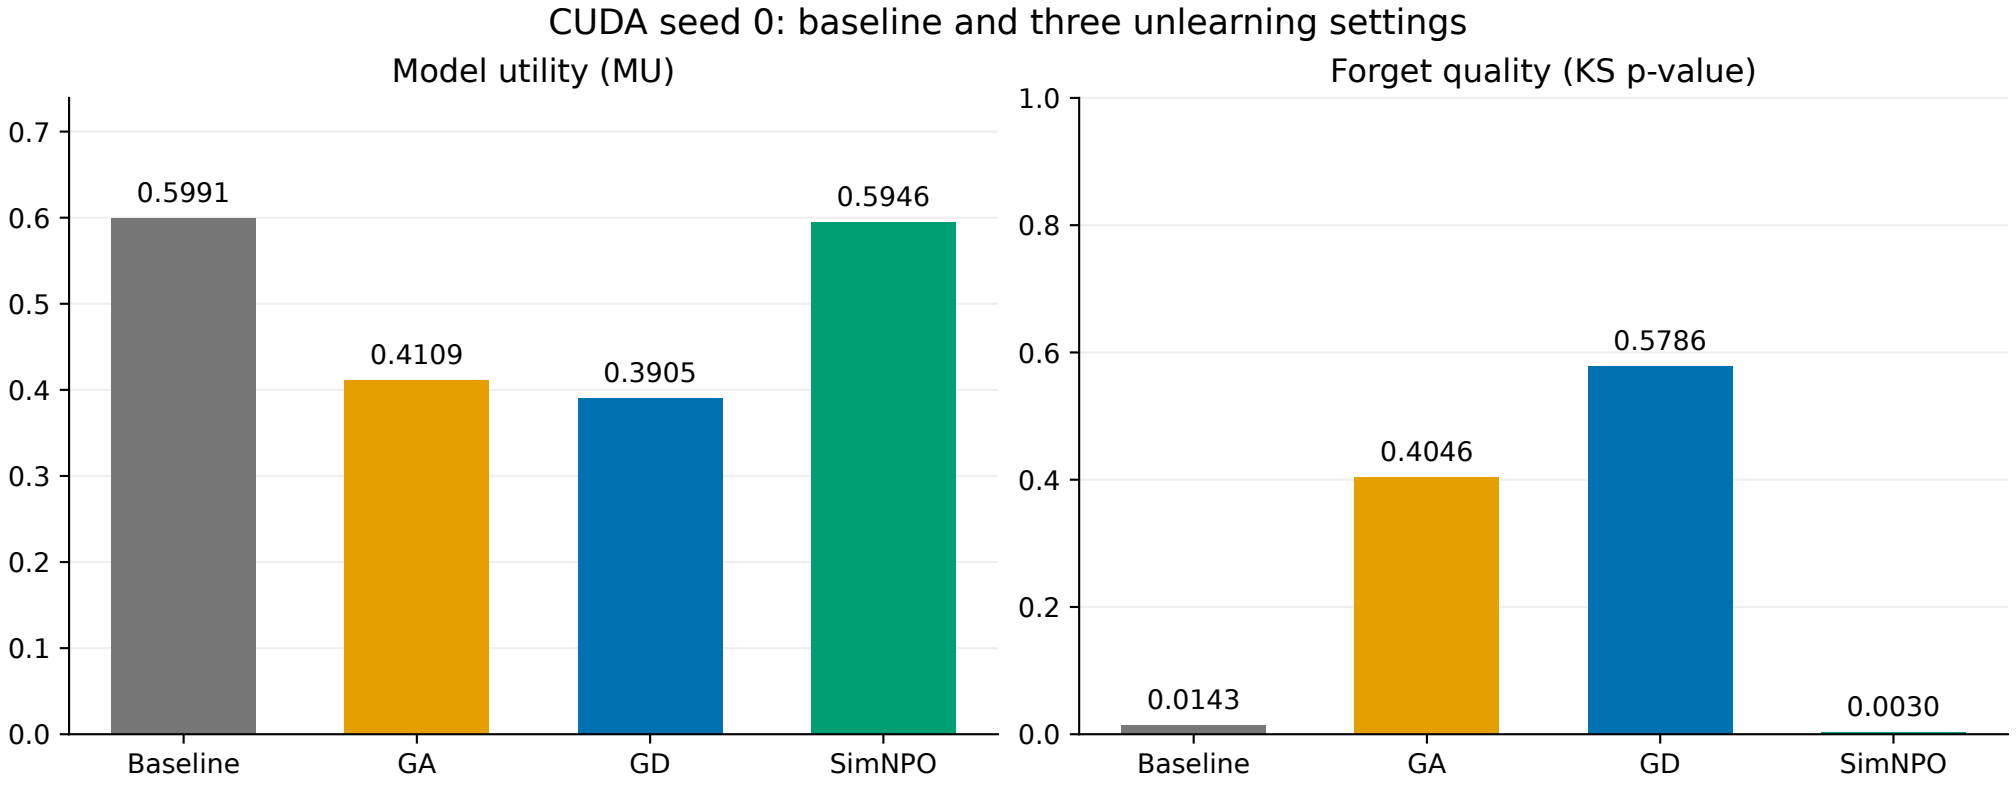
CUDA seed 0: baseline and three unlearning settings
Model utility (MU)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.5991
0.4109
0.3905
0.5946
Baseline
GA
GD
SimNPO
Forget quality (KS p-value)
0.0
0.2
0.4
0.6
0.8
1.0
0.0143
0.4046
0.5786
0.0030
Baseline
GA
GD
SimNPO
Trained methods: 10 updates / 320 visits; one run per setting, no error bars